



- | **AI Project Management Agent: From Planning to Execution**
- | **Subtitle:**
- | **Using Preference Fine-Tuning (PFT) + Reinforcement Fine-**
- | **KeyiMe (Rag) with Jira, GitHub, Notion, Slack**
- | **Turn LLMs into senior PM copilots that can plan, track, and communicate across your real tools.**



- Fragmented across Jira, GitHub, Notion, Slack
- Full of repetitive status updates and manual planning
- Often inconsistent in quality between teams

| LLMs today:

- Can "talk about" projects,
- But aren't reliably grounded in your tickets, repos, docs, and chats.

| Opportunity:

An agent that can understand the project, propose realistic plans, and operate tools safely.

1. Enterprise-grade planning



- Hierarchical plans, realistic timelines, clear risks
- Tool-native execution
- Create/update Jira issues, GitHub PR summaries, Notion specs, Slack updates
- Reliable & auditable
- Human-in-the-loop approvals
- Policy guardrails and logs
- Composable architecture
- Swap model, tools, or data sources without redoing everything

Four main layers:



- Foundation & Alignment Layer
- Base LLM + SFT + PFT (DPO) + safety policies
- Reasoning & Planning Layer
 - Project planner, timeline estimator, risk engine, dependency resolver
- Tool & Data Layer
 - Jira, GitHub, Notion, Slack connectors
 - Vector search over project docs, tickets, specs
- Orchestration & UX Layer
 - Agent controller, workflows, approvals
- UI: chat, side-panel in Jira/Notion, Slack app

- Start with a strong base LLM (general + coding).



- SFT on:
- PRDs, sprint docs, design docs, runbooks
- PFT / DPO on preference datasets:
- Good vs bad project breakdowns
- Strong vs weak timelines
- Good vs bad risk registers
- High vs low quality stakeholder comms

| Outcome:

| The model "thinks like a senior PM."

- Offline environment that simulates:

- Jira: create/update tickets, change status, assign
- GitHub: read PR metadata, post summary comments
- Notion: read/write pages
- Slack: draft & send updates (with approval)
- Reward examples:
 - +1 if the right Jira field is set
 - +1 if the right repo/branch is referenced
 - +1 if Slack summary matches actual issue state
 - -1 if invalid API calls or inconsistent states

| **Outcome:**

- LLM Core (PFT + RFT trained)



- Agent Orchestrator
- Decides when to call tools vs. reason internally
- Memory & Retrieval
- Vector DB over specs, tickets, docs
- Tool Adapters
- JiraClient, GitHubClient, NotionClient, SlackClient
- Policy & Safety Layer
- Access scope, rate limiting, approval policy

1. User asks: "Help me plan Q3 release for Feature X."

- Agent:
 - Retrieves relevant docs (Notion/Confluence)
 - Fetches Jira epics + open bugs
 - Checks GitHub branches/PRs
 - Fetches last Slack status updates
- LLM:
 - Synthesizes a project plan, timelines, risks
 - Proposes Jira changes & Slack message drafts
- Human:
 - Reviews & approves changes / messages



"Create a 3-sprint plan to ship Feature X."

Agent:

- Gets existing Jira epics/issues
- Groups into sprints by priority, dependencies, team capacity
- Proposes new tickets for gaps
- Writes a Slack summary for the team

Human approves:

- Agent creates/updates Jira issues
- Posts Slack summary draft to a channel



| **"List top 10 risks for Feature X and mitigation plans."**

| **Agent:**

- Reads spec, Jira issues, known blockers
- Uses risk-register patterns to propose:
- Risk name, impact, probability, owner, mitigation
- Drafts a Notion "Risk Register" section

| **Human approves and edits as needed.**



- Human rating of plans (clarity, realism, coverage)
- Tool correctness
- % of successful Jira/GitHub/Notion actions
- Latency & UX
- Time from request to draft plan
- Adoption & satisfaction
- PM/EM satisfaction scores
- Reduction in manual PM work

- Scoped permissions per tool



- All actions logged with:
- Who requested
- What the agent proposed
- What was executed
- Sensitive operations require:
- Approval
- Or "draft-only" mode

- Phase 1: Read-only advisor (no writes, only suggestions)



- Phase 2: Draft creator (Jira/Slack drafts, Notion pages)
- Phase 3: Limited auto-execution with guardrails
- Phase 4: Org-wide deployment + custom policies & KPIs

End of Slide Deck

